



Fast Fourier Transform





Overview

- ◇ What does it do and why do we use it?
- ◇ Implementation in C using MPI
- ◇ Performance





Mechanics and Usage

Why do we care?



Fourier Transform

- ◇ In essence, it allows us to move from one domain to another
- ◇ Creates what is called a “Fourier Series” from an expression
- ◇ This series is representative of the original expression and can be reversed
- ◇ For our case, we will use it go decompose complex signals into parts for analysis






Discretization

- ◇ In the real world, we do not have expressions to represent signals
- ◇ Given a set of samples, we can run operations on them to achieve the same sort of results.
- ◇ This process is of $O(n^2)$ in a naive DFT

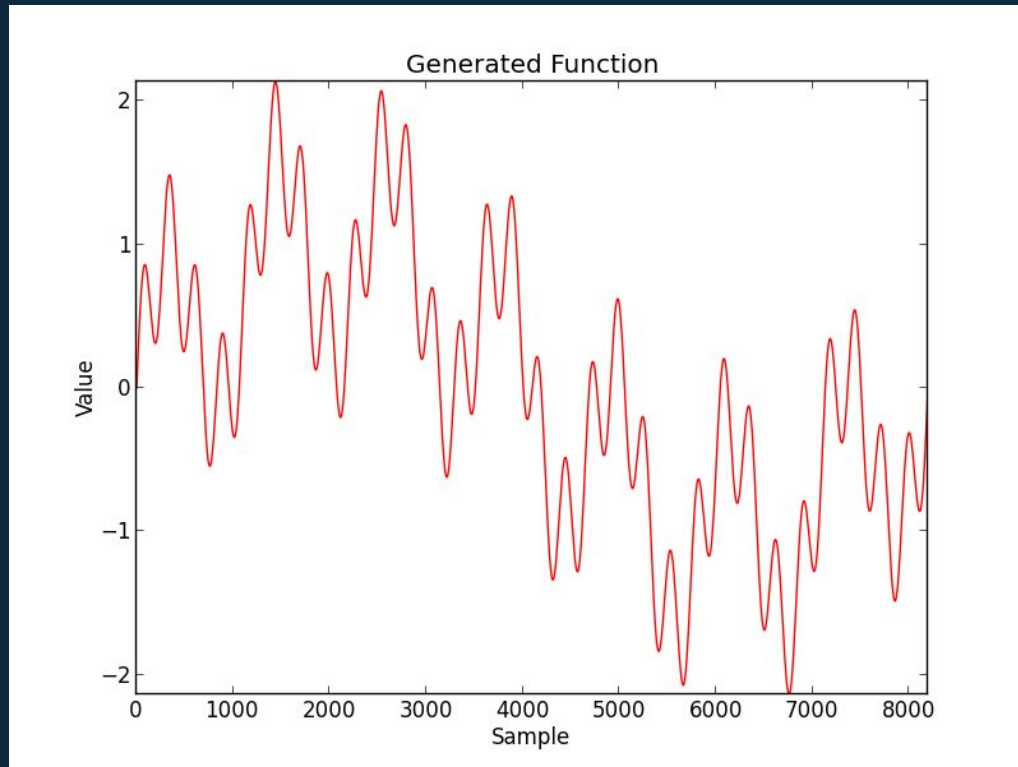



$$X_k = \frac{1}{N} \sum_{n=0}^{N-1} x_n e^{i2\pi k \frac{n}{N}}$$


To find the energy at a particular frequency, spin your signal around a circle at that frequency and average a bunch of points along that path




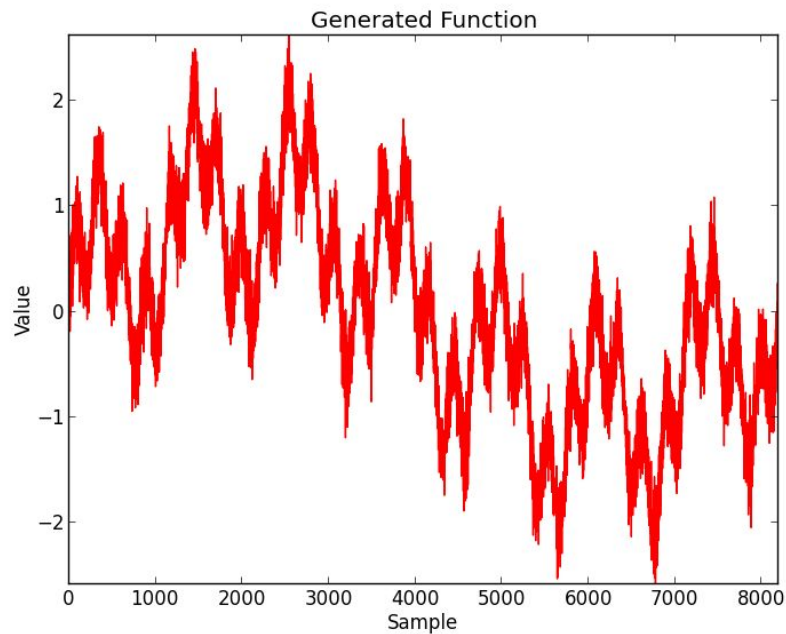
Time Domain Signal



Frequency Domain Signal



Noise Resistance





Fast Fourier Transform

- ◇ Most commonly used FFT is the Cooley - Tukey Algorithm
- ◇ Process is expressed much more succinctly and requires fewer operations
- ◇ Is inherently recursive, which can be tricky with large problem sizes
- ◇ Reduces work to $O(n * \log n)$ for 1D FFT
- ◇ A 2D FFT can be achieved by applying the operation over all rows, then all columns



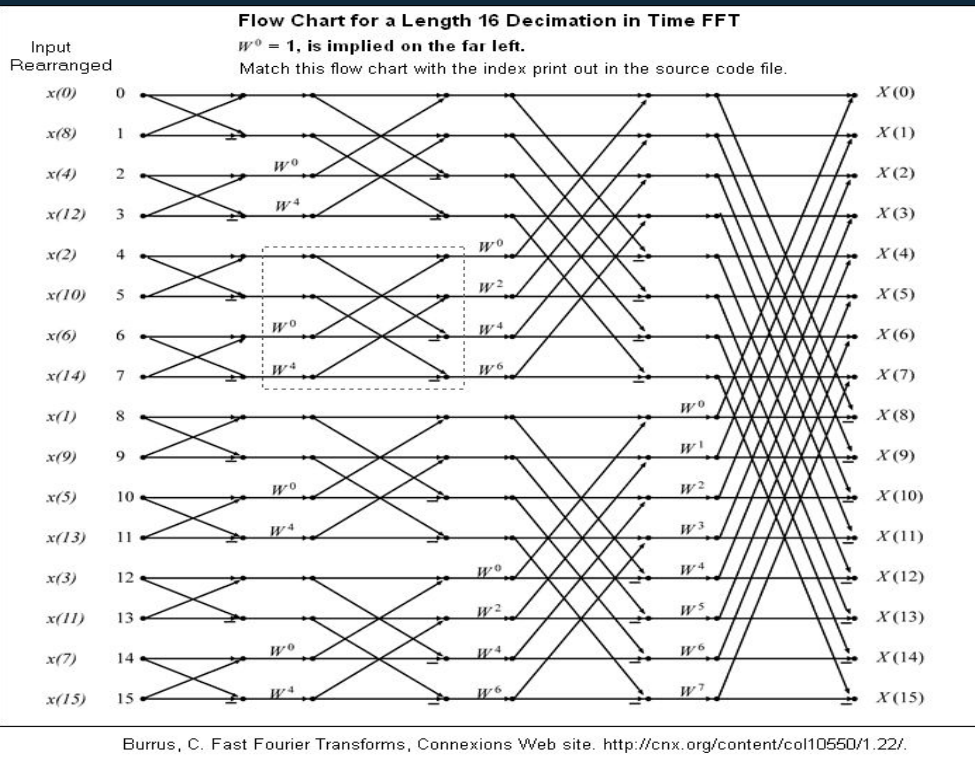


Implementation

Serial, OpenMP, and MPI



Butterfly Operations





Recursive Algorithm

$X_{0,\dots,N-1} \leftarrow \text{ditfft2}(x, N, s):$

if $N = 1$ then

$X_0 \leftarrow x_0$

else

$X_{0,\dots,N/2-1} \leftarrow \text{ditfft2}(x, N/2, 2s)$

$X_{N/2,\dots,N-1} \leftarrow \text{ditfft2}(x+s, N/2, 2s)$

for $k = 0$ to $N/2-1$

$t \leftarrow X_k$

$X_k \leftarrow t + \exp(-2\pi i k/N) X_{k+N/2}$

$X_{k+N/2} \leftarrow t - \exp(-2\pi i k/N) X_{k+N/2}$

endfor

endif

DFT of $(x_0, x_s, x_{2s}, \dots, x_{(N-1)s})$:

trivial size-1 DFT base case

DFT of $(x_0, x_{2s}, x_{4s}, \dots)$

DFT of $(x_s, x_{s+2s}, x_{s+4s}, \dots)$

combine DFTs of two halves





General Ideas

- ◇ We can make our operations easier by sorting in bit reverse order
- ◇ This way, we can use an iterative approach, which solves several problems regarding stack and memory
- ◇ Requires problem size of 2^K nodes
- ◇ Allows for more parallelizable functions later on





OpenMP

- ◇ I quickly discovered that this problem benefited very little from parallel optimizations in 1D
- ◇ I turned my focus to the 2D problem
- ◇ Decomposed problem into “Stripes” of multiple rows of whole
- ◇ Ran FFT over each item in chunk, transposed information and repeated
- ◇ Still did not see much of a speed up in the OpenMP version





MPI

- ◇ Approach was very similar to OpenMP version
- ◇ Communication was the most important element of my changes
- ◇ Made all operations run based on 1D complex arrays (no complicated structures)
- ◇ Special operations were needed to package messages for distribution



Row Phase

Proc 0	1	2	3	4	5	6	7	8
	9	10	11	12	13	14	15	16
	17	18	19	20	21	22	23	24
	25	26	27	28	29	30	31	32
Proc 1	33	34	35	36	37	38	39	40
	41	42	43	44	45	46	47	48
	49	50	51	52	53	54	55	56
	57	58	59	60	61	62	63	64



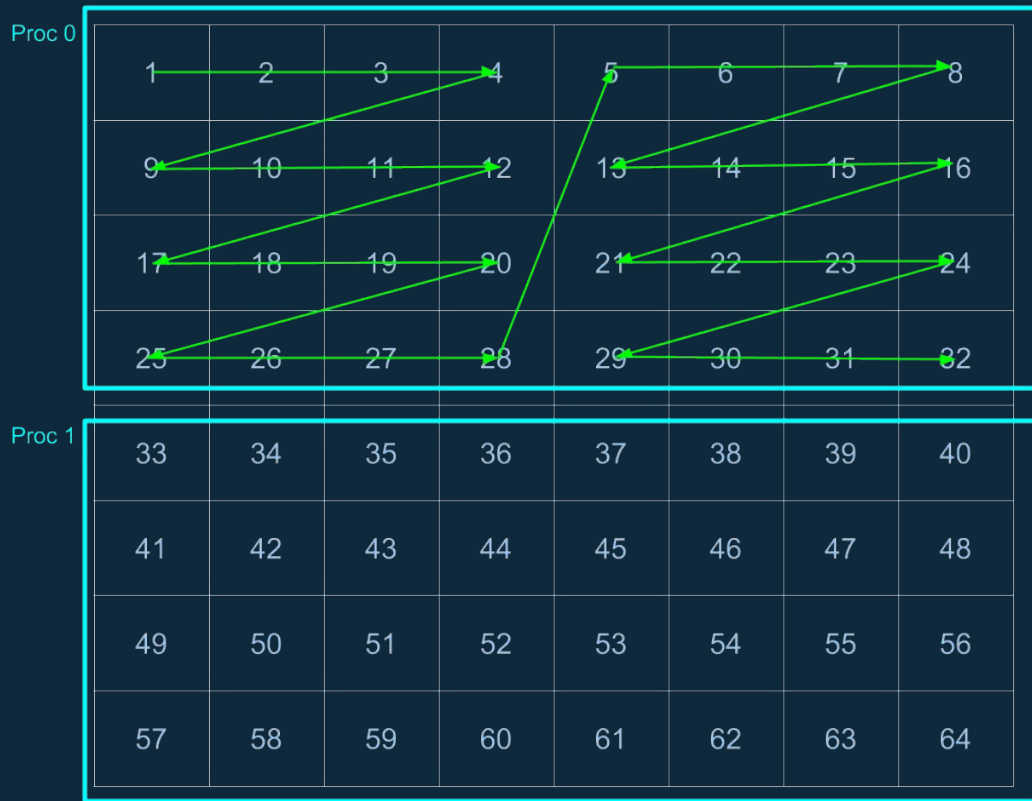
Column Phase

Proc 0	1	2	3	4	5	6	7	8
	9	10	11	12	13	14	15	16
	17	18	19	20	21	22	23	24
	25	26	27	28	29	30	31	32
	33	34	35	36	37	38	39	40
	41	42	43	44	45	46	47	48
	49	50	51	52	53	54	55	56
	57	58	59	60	61	62	63	64
Proc 1								



Data read pattern for
transfer into send
buffer

All to all would cut
data along jump from
28 to 5 in this case





Parallel Performance

Large problems, fast times




Maneframe Strong Scaling

N	N * N	Processors	Time	Speedup (real)	Efficiency (real)
4096	16777216	1	9.712320089	1	1
4096	16777216	2	5.76313591	1.685249184	0.842624592
4096	16777216	4	3.248805046	2.989505357	0.7473763393
4096	16777216	8	1.806185007	5.377256511	0.6721570639
4096	16777216	16	1.490156889	6.517649357	0.4073530848
4096	16777216	32	1.290680885	7.524958493	0.2351549529
4096	16777216	64	1.214707851	7.99560164	0.1249312756




Run Time (4096 * 4096)




Strong Scaling (4096 * 4096)

Speedup (real) Efficiency (real)



Weak Scaling Efficiency (Work per Proc = 65536)





Notes on Performance

- ◇ Although uncertain, I believe that network topology has a great deal of impact
- ◇ Weak scaling was poor, this is due to the fact that communication scales with size
- ◇ Due to limitations of size, only powers of two processors can be tested
- ◇ Was already very fast in serial, so all speedup is dependent on communications





Final Thoughts

Considerations and Future Work



Moving Forward

- ◇ Remove dependence on square matrix inputs
- ◇ Add real time system functionality, make callable from other applications
- ◇ Interface with image libraries for image signal analysis
- ◇ Run tests on higher numbers of procs, larger files
- ◇ Look into ways of distributing file I/O for faster load and dump operations





Thank You!





References

- [1]"An Interactive Guide To The Fourier Transform – BetterExplained", *Betterexplained.com*, 2017. [Online]. Available: <https://betterexplained.com/articles/an-interactive-guide-to-the-fourier-transform/>. [Accessed: 08- May- 2017].
- [2]"Butterfly Diagram", *Iowahills.com*, 2017. [Online]. Available: <http://www.iowahills.com/Example%20Code/FFT%20Butterfly%20Diagram.png>. [Accessed: 11- May- 2017].
- [3]"Implementing FFTs in Practice", *Cnx.org*, 2017. [Online]. Available: <http://cnx.org/contents/ulXtQbN7@15/Implementing-FFTs-in-Practice>. [Accessed: 11- May- 2017].
- 